



House Price Prediction

An end-to-end MLOps pipeline for real estate price prediction



Data Processing



Model Training



Containerization



Deployment

Project Foundation & Dataset

Dataset

 1,000+

House Price Records

Comprehensive property dataset with multi-dimensional features for accurate price prediction.

Features

-  Square Footage
-  Bedrooms
-  Bathrooms
-  Location
-  Year Built
-  Condition

Price Statistics

Mean:  \$553,234

Median:  \$495,000

Range:  \$249K–\$1.25M



Technology Stack

CI/CD Automation

GitHub Actions automates data processing and container publishing on every Git push.

Model Packaging

FastAPI + Uvicorn converts trained models into production-ready web APIs.



Monitoring Stack

Prometheus + Grafana provides real-time performance insights and alerting.

User Interface

Streamlit provides the user-friendly web application for predictions.

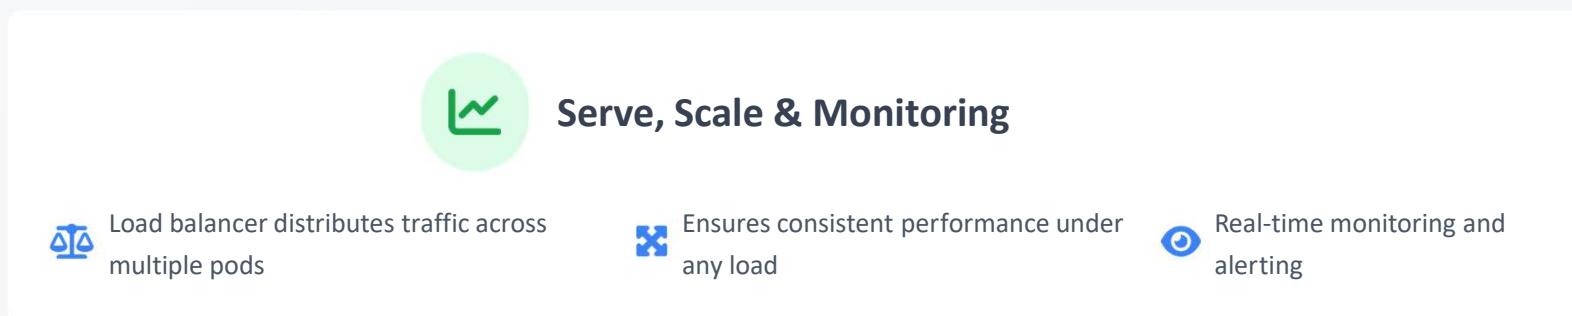
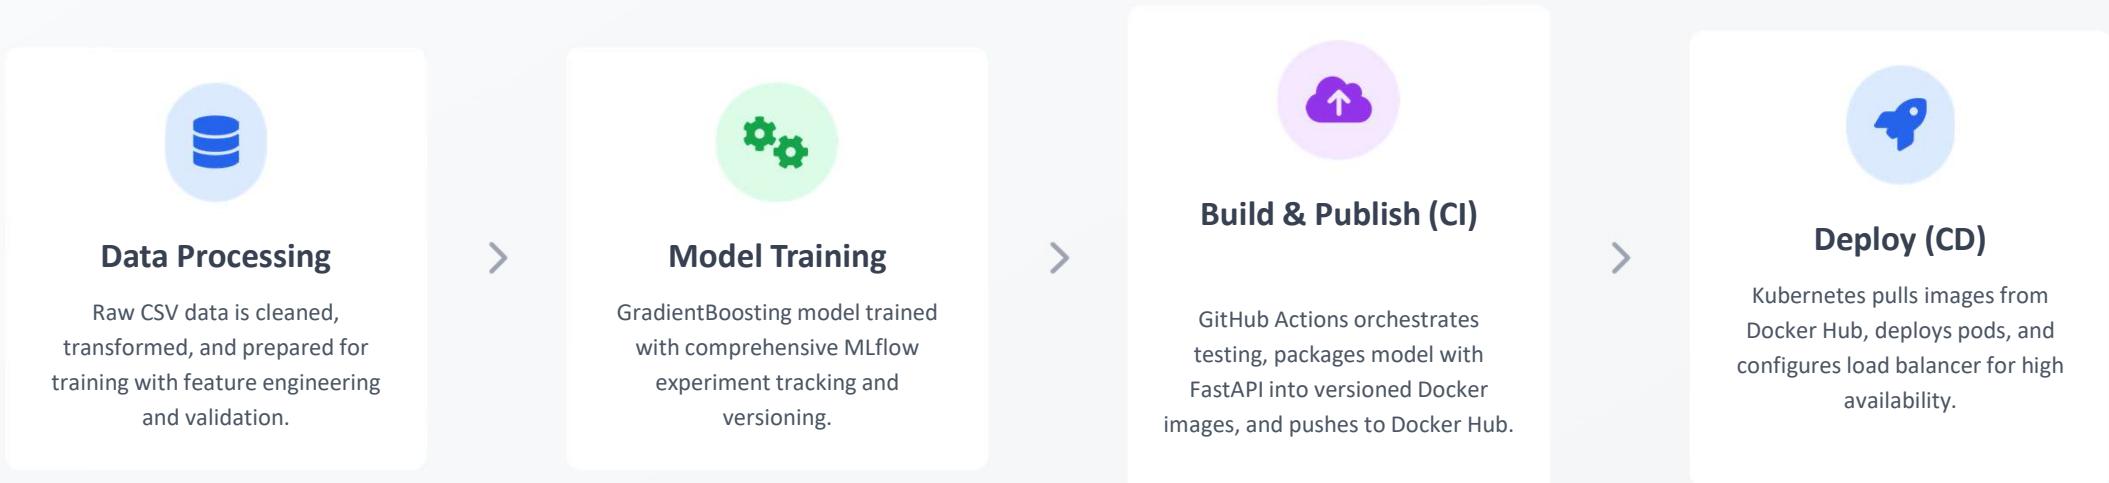
Container Registry

DockerHub stores versioned container images for consistent deployment.

Model Deployment

Kubernetes (kind) manages and scales the containerized application.

End-to-End Workflow Overview



GitHub Actions

Docker

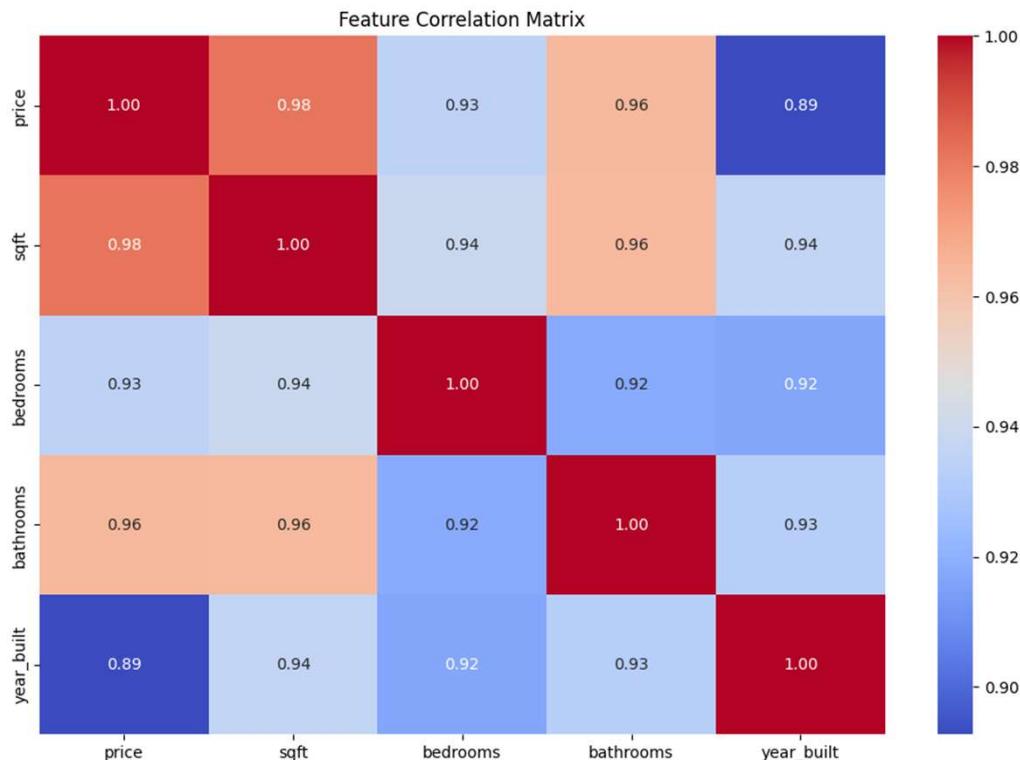
Kubernetes

FastAPI

Prometheus + Grafana

Feature Relationships & Analysis

Feature Correlation Matrix



Key Insights

Strongest Correlations

Price shows strong positive correlation (>0.98) with square footage, indicating size is a primary pricing factor.

Property Size Impact

Square footage demonstrates the strongest linear relationship with price among all features.

Amenity Value

Bathrooms show significant correlation with price(0.96), highlighting the importance of luxury amenities.

Model Performance Comparison

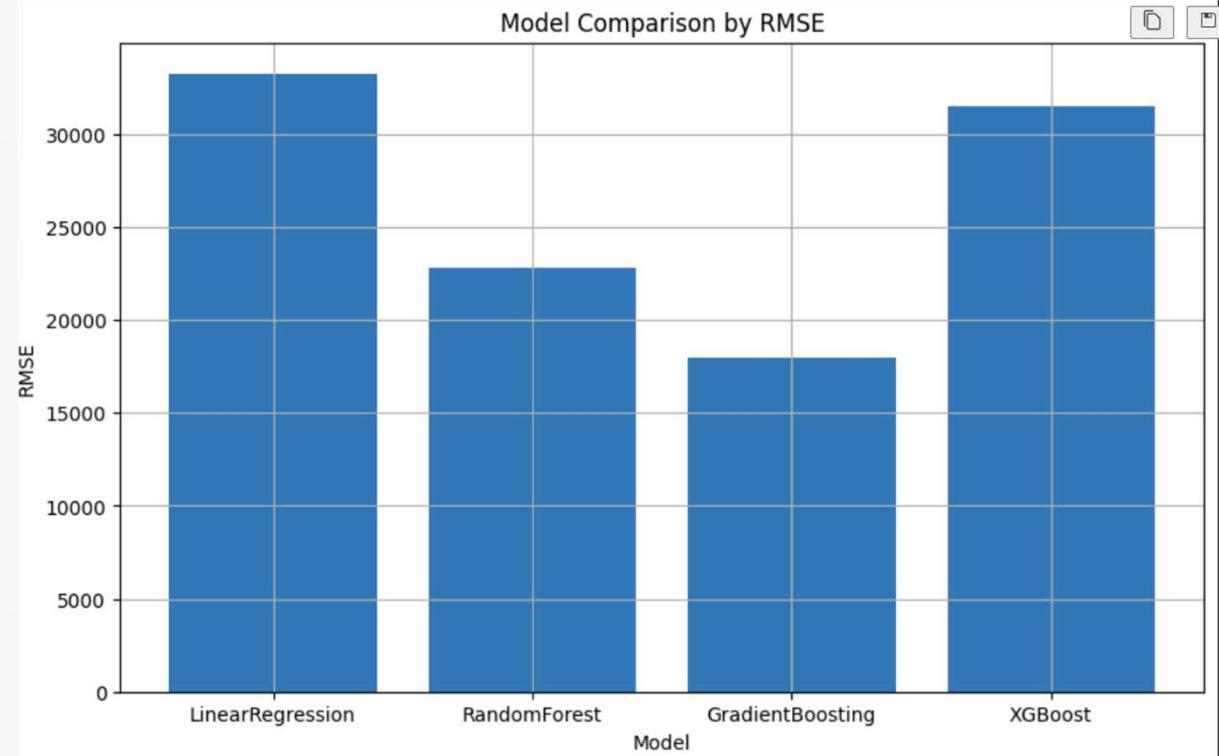
Comparison of different ML algorithms for house price prediction. GradientBoosting model shows optimal performance for this regression task.

GradientBoosting

R ² Score:	0.99
RMSE (USD):	\$17,968
Error %	3.25%

Performance Characteristics

Model	R ² Score	RMSE (USD)	Error %
GradientBoosting	0.9931	\$17,968	3.25%
Random Forest	0.9888	\$22,818	4.12%
XGBoost	0.9787	\$31,513	5.70%
Linear Regression	0.9762	\$33,267	6.01%



MLFlow for Experiment Tracking

Overview

Comprehensive Experiment Management

MLflow provides comprehensive tracking and versioning for all model experiments, ensuring reproducibility and easy comparison.

-  Reproducibility
-  Easy Comparison
-  Version Control

Key Features



Hyperparameter Tracking

Record and compare model parameters across experiments



Metrics Logging

Track performance metrics to identify optimal models



Artifact Storage

Save model artifacts and configurations



Configuration Export

Export to model_config.yaml for deployment



Model Export

Export best_model.pkl for production use

Workflow

Configure Parameters

Set estimators to 250 for optimal performance

01

Track Experiments

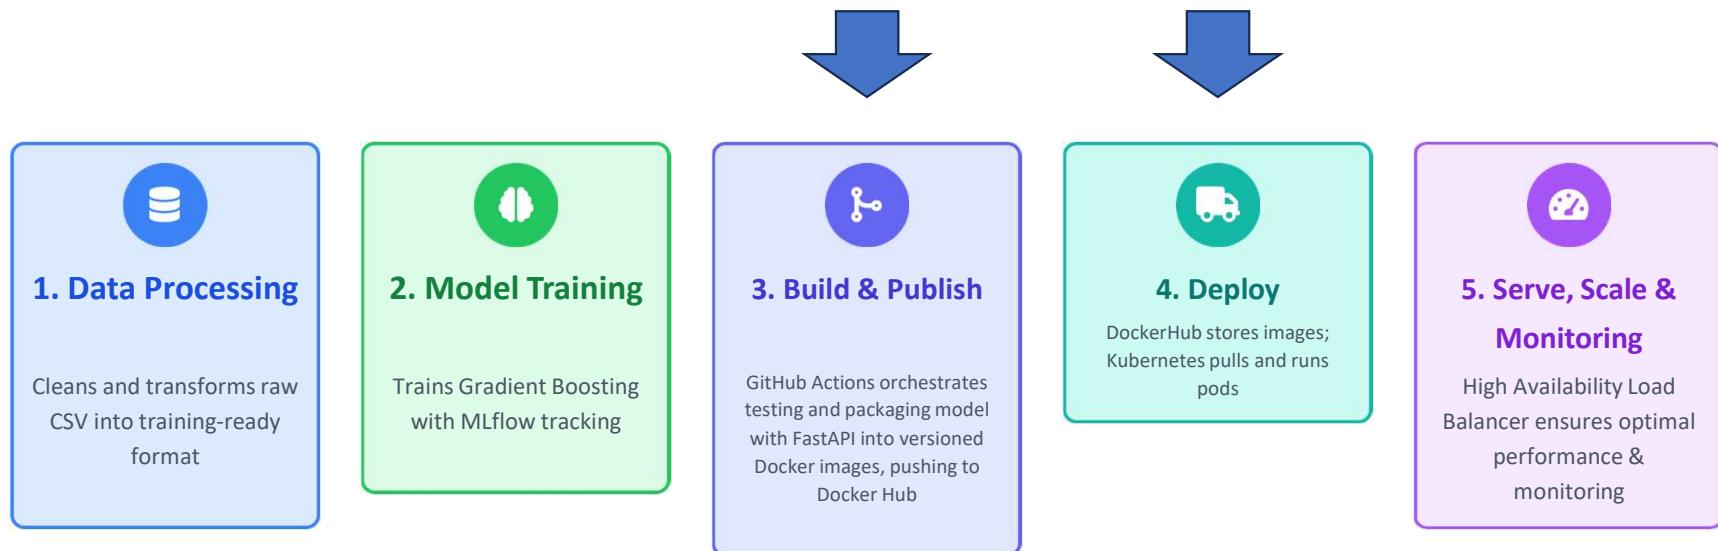
Log all metrics and model versions

02

Save Artifacts

Export configuration and trained model

Workflow Overview

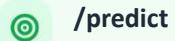


Production API & User Interface



FastAPI Backend

Production-ready API for house price prediction:



/predict

Single prediction endpoint



/batch-predict

Bulk processing capability



/metrics

Performance monitoring



Streamlit Frontend

User-friendly interface for predictions:



Square Footage



Bedrooms



Bathrooms



Location



Prediction Results

\$450K - \$650K

Based on 100 similar properties

Streamlit Application

User Interface

Streamlit provides a web app for house price predictions with:

- ✓ Simple input forms for property details
- ✓ Real-time price estimation
- ✓ Visualization of influencing factors

Key Input Fields

 Square Footage

 Bedrooms & Bathrooms

 Location

Prediction Results

House Price Prediction

A simple MLOps demonstration project for real-time house price prediction

Square Footage:

Bedrooms

3

Bathrooms

2

Location

Suburban

Year Built:

Predict Price

Prediction Results

Connecting to API at: <http://model:8000/predict>

\$300,826

Confidence Score

92%

Model Used

XGBoost

Price Range

\$270,743.4 - \$330,908.6

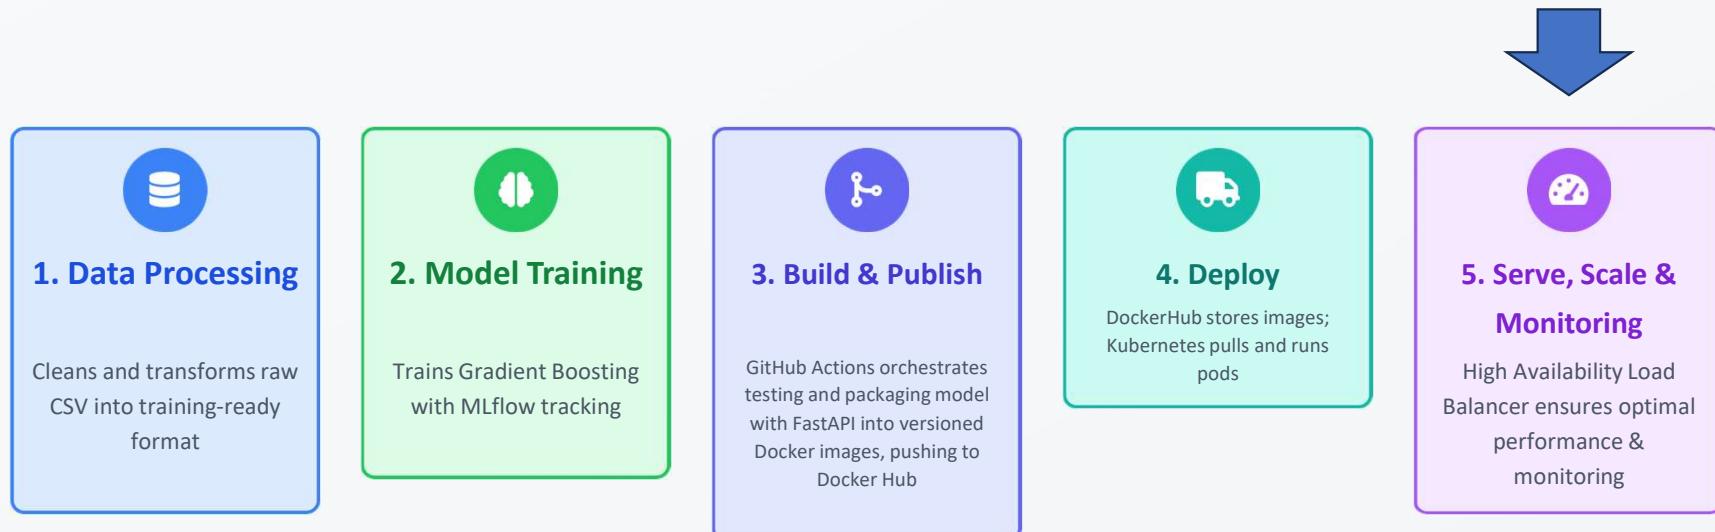
Prediction Time

0.12 seconds

Top Factors Affecting Price:

- Square Footage
- Number of Bedrooms/Bathrooms

Workflow Overview



Docker Containerization

The Magic Packaging Box

Docker is like a sturdy, special box that securely packs your application with everything it needs to run.

This "box" ensures your application runs consistently across any environment, solving the "But it worked on my machine!" problem.



Universal Compatibility

Works identically on any computer—no more environment-specific issues.

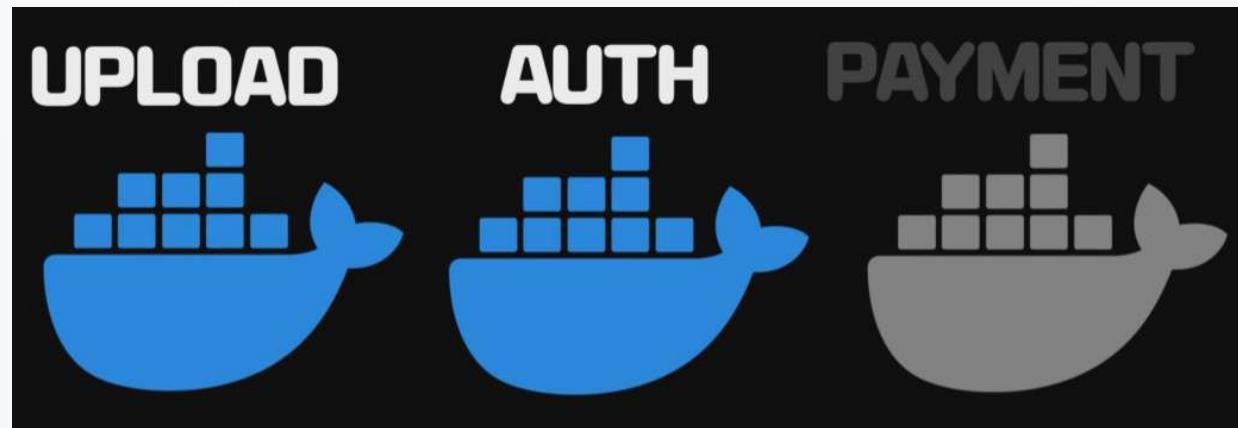
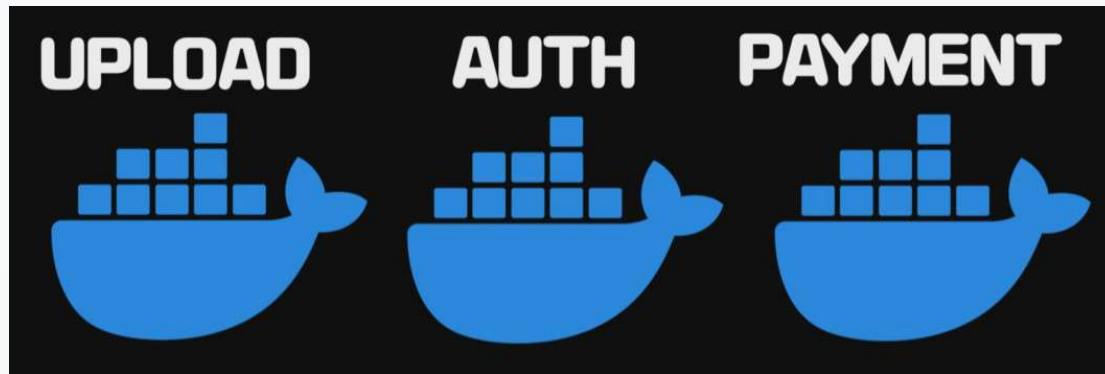


Complete Package

Packs your application, dependencies, and configuration into one perfect container.



Why Kubernetes?



Kubernetes Orchestration

Kubernetes orchestrates Docker containers like a factory manager overseeing production at scale.



Auto-Scaling

Automatically adds containers during peak traffic and removes them during quiet periods.



Load Distribution/Orchestration

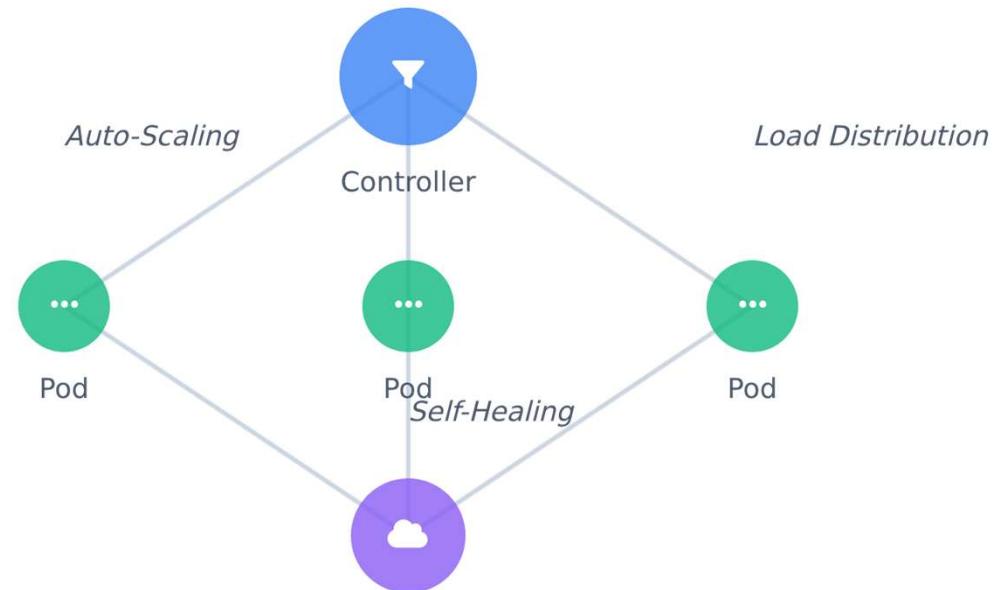
Distributes containers evenly across many servers during high traffic.



Self-Healing

Detects failures instantly and moves containers to healthy machines.

Orchestration Visualization



Kubernetes Core Resources



Pod

- The smallest execution unit Kubernetes manages
- A collection of one or more Docker containers
- Share networking and storage resources
- Example: FastAPI model or Streamlit UI



Deployment

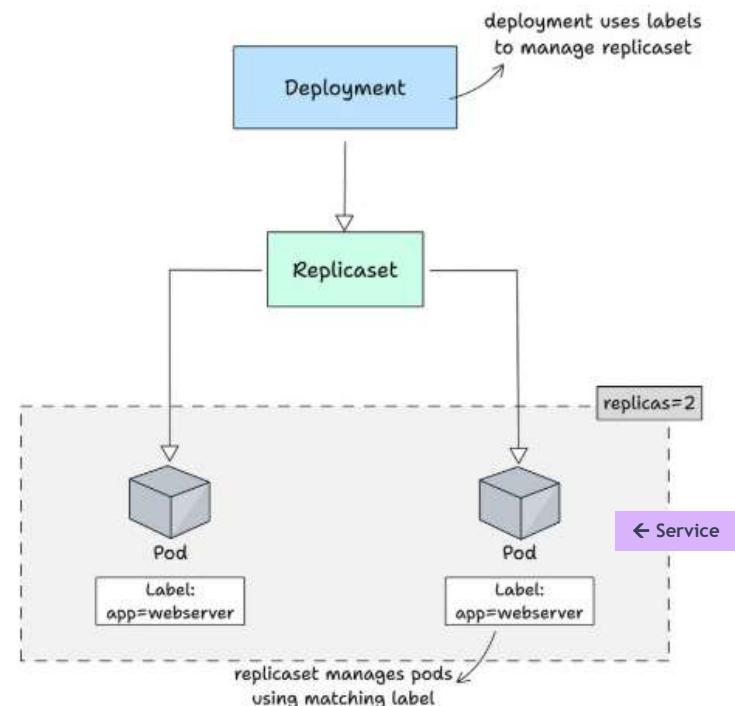
- Management rule that automatically controls and replicates Pods
- Ensures FastAPI Model Pods are kept at the specified number of replicas
- Automatically recovers them if one fails



Service

- Network administrator that provides a fixed access point
- Handles traffic distribution (Load Balancing)
- Ensures Streamlit UI can talk to the Model Pods reliably

Kubernetes Architecture



● Pod

● Deployment

● Service

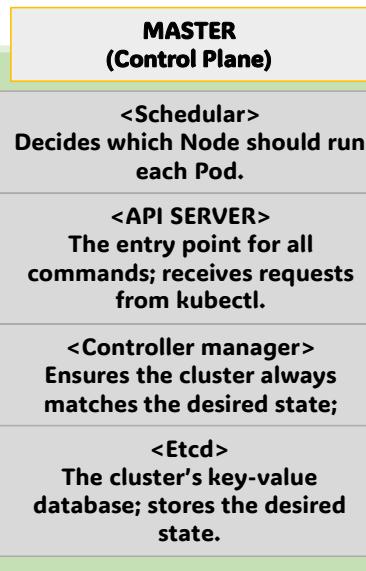
Production Deployment Status

Cluster Overview



Kubectl

YAML File



Container Registry:

DockerHub

Worker Node 1

Model container
Model container
Streamlit container
Streamlit container

Runtime
Kube-proxy
Kubelet

...Worker Node '2'

...Worker Node 'n'...

Underlying infrastructure



Physical



Virtual



Private



Public



Hybrid

```
honor@LAPTOP-NG18QTQ5 MINGW64 ~/spicedAcademy/mlops/bootcamp/mlops_house-price-predictor (main)
$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
kube-ops-view-7d756f7b94-jxpq6   1/1    Running   6 (8h ago)  12d
model1-6656fc88c7-dd485        1/1    Running   6 (8h ago)  11d
model1-6656fc88c7-h9162        1/1    Running   6 (8h ago)  11d
streamlit-cb664c6d7-w4ldn       1/1    Running   6 (8h ago)  11d
streamlit-cb664c6d7-x9n8c       1/1    Running   6 (8h ago)  11d
```

FastAPI Model Server

Two model replicas are running, successfully providing the predictions service.

- ✓ Fulfils requirements for load balancing and high availability

Streamlit UI

Two UI replicas are running, enabling the application to distribute user experience effectively across instances.

- ✓ Ensures consistent user experience

Resources and Reliability in Action

- Load Distribution(Service): Spreads incoming traffic evenly across 2 Pods.
- Self-Healing(by Controller Manager): If FastAPI Pod #1 crashes, Kubernetes automatically creates a new Pod.
- Autoscaling(yaml): During a traffic spike, Pods scale from 2 to 4 automatically.

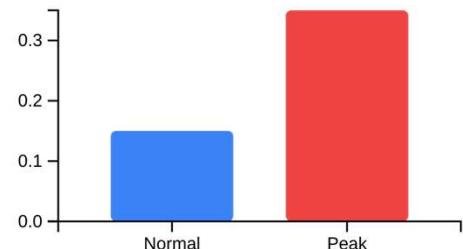


Grafana Monitoring Dashboard

Real-time Performance Insights

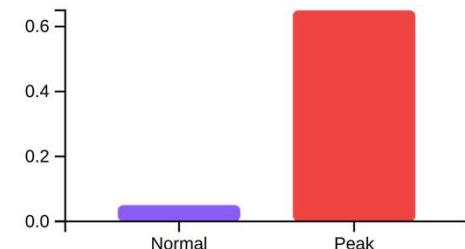
Prometheus + Grafana provides real-time insights into our House Price Prediction service, enabling us to identify performance bottlenecks.

Request Rate



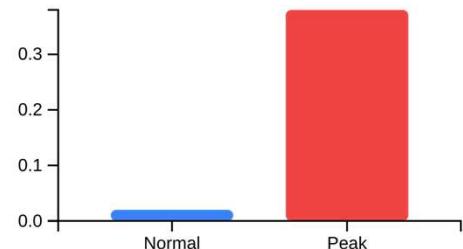
/metrics: No change (0.1–0.2)
/predict: Surge (0.1 to ~0.4)

Latency (95th Percentile)

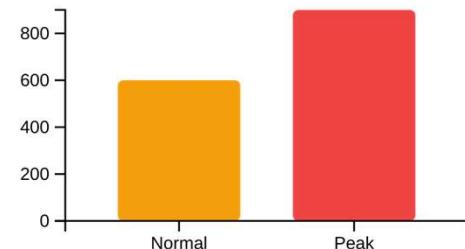


/metrics: No change (~0)
/predict: Surge (0.4 to 0.9)

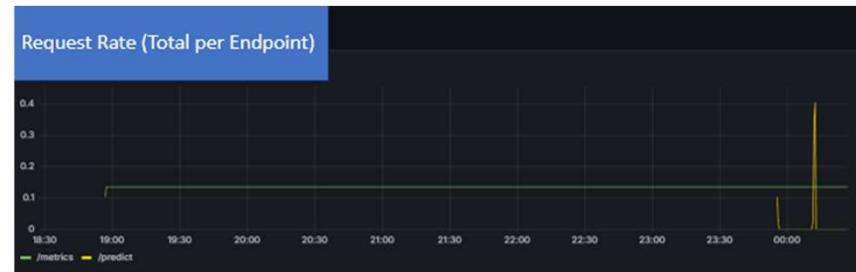
Request Size



Response Size



Grafana Dashboard



Peak Load Analysis

- ↗ A sudden influx of requests, increasing system load
- ☒ Response time increased to nearly 1 second
- ↖ Response data size increased by approximately 50%

During peak activity, the /predict endpoint experiences:
latency failure during traffic spikes